

In [1]:

```
import pandas as pd
import os
import numpy as np
import seaborn as sns
import joblib as jb
import sklearn
import pydotplus
import matplotlib.pyplot as plt,pydotplus
```

In [2]:

```
from sklearn.preprocessing import LabelEncoder
```

In [3]:

```
pd.read_excel('Combined_Updated.xlsx')
```

Out[3]:

| | AgeCategory | Workclass | Education | EducationNum | MaritalStatus | Occupation | Relationship | Sex | NativeCountry | Race | FnlwgtC: |
|-------|-------------|-----------|-----------|--------------|---------------|------------|--------------|-----|---------------|------|----------|
| 0 | 3 | 6 | 9 | 12 | 4 | 0 | 1 | 1 | 38 | 4 | |
| 1 | 0 | 5 | 9 | 12 | 2 | 3 | 0 | 1 | 38 | 4 | |
| 2 | 3 | 3 | 11 | 8 | 0 | 5 | 1 | 1 | 38 | 4 | |
| 3 | 0 | 3 | 1 | 6 | 2 | 5 | 0 | 1 | 38 | 2 | |
| 4 | 3 | 3 | 9 | 12 | 2 | 9 | 5 | 0 | 4 | 2 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 48837 | 3 | 3 | 9 | 12 | 0 | 9 | 1 | 0 | 38 | 4 | |
| 48838 | 2 | 3 | 11 | 8 | 6 | 9 | 2 | 1 | 38 | 2 | |
| 48839 | 3 | 3 | 9 | 12 | 2 | 9 | 0 | 1 | 38 | 4 | |
| 48840 | 0 | 3 | 9 | 12 | 0 | 0 | 3 | 1 | 38 | 1 | |
| 48841 | 3 | 4 | 9 | 12 | 2 | 3 | 0 | 1 | 38 | 4 | |

48842 rows × 15 columns



In [4]:

```
data=pd.read_excel('Combined_Updated.xlsx')
```

In [5]:

```
data.describe()
```

Out[5]:

| | AgeCategory | Workclass | Education | EducationNum | MaritalStatus | Occupation | Relationship | Sex | NativeCountr |
|-------|--------------|--------------|--------------|--------------|---------------|--------------|--------------|--------------|--------------|
| count | 48842.000000 | 48842.000000 | 48842.000000 | 48842.000000 | 48842.000000 | 48842.000000 | 48842.000000 | 48842.000000 | 48842.000000 |
| mean | 1.757893 | 3.099668 | 10.288420 | 9.078089 | 2.618750 | 6.152819 | 1.443287 | 0.668482 | 36.43366 |
| std | 1.373868 | 1.110810 | 3.874492 | 2.570973 | 1.507703 | 3.968837 | 1.602151 | 0.470764 | 6.03153 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.00000 |
| 25% | 0.000000 | 3.000000 | 9.000000 | 8.000000 | 2.000000 | 3.000000 | 0.000000 | 0.000000 | 38.00000 |
| 50% | 3.000000 | 3.000000 | 11.000000 | 9.000000 | 2.000000 | 7.000000 | 1.000000 | 1.000000 | 38.00000 |
| 75% | 3.000000 | 3.000000 | 12.000000 | 11.000000 | 4.000000 | 9.000000 | 3.000000 | 1.000000 | 38.00000 |
| max | 3.000000 | 7.000000 | 15.000000 | 15.000000 | 6.000000 | 13.000000 | 5.000000 | 1.000000 | 40.00000 |



In [6]:

```
data["Class"].value_counts()
```

Out[6]:

```
0    37155
1    11687
Name: Class, dtype: int64
```

In [9]:

```
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
xc=['AgeCategory', 'Workclass', 'Education', 'EducationNum', 'MaritalStatus', 'Occupation', 'Relationship', 'NativeCountry', 'Sex',
y=['Yes', 'No']
all_input=data[xc]
all_class=data['Class']
```

In [10]:

```
(x_train,x_test,y_train,y_test)=train_test_split(all_input,all_class,train_size=0.67,random_state=10)
```

In [11]:

```
from sklearn.ensemble import RandomForestClassifier
clf=RandomForestClassifier(criterion="gini",max_depth=8,random_state=5)
```

In [13]:

```
clf.fit(x_train, y_train)
```

Out[13]:

```
RandomForestClassifier
RandomForestClassifier(max_depth=8, random_state=5)
```

In [21]:

```
Y_train_pred=clf.predict(x_train)
Y_test_pred=clf.predict(x_test)
```

In [17]:

```
from sklearn.metrics import accuracy_score
accuracy_score(y_train,Y_train_pred)
```

Out[17]:

```
0.8458318054027625
```

In [19]:

```
from sklearn.model_selection import cross_val_score
cross_val_score(clf,x_train, y_train,cv=10)
```

Out[19]:

```
array([0.83195845, 0.84631836, 0.8395967 , 0.83226398, 0.83251834,
       0.84749389, 0.83832518, 0.8358802 , 0.84107579, 0.84535452])
```

In [23]:

```
from sklearn.metrics import classification_report
print(classification_report(Y_test_pred,y_test))
```


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.94 | 0.86 | 0.90 | 13331 |
| 1 | 0.52 | 0.72 | 0.61 | 2787 |
| accuracy | | | 0.84 | 16118 |
| macro avg | 0.73 | 0.79 | 0.75 | 16118 |
| weighted avg | 0.87 | 0.84 | 0.85 | 16118 |

In [25]:

```
features = data.columns
importances=clf.feature_importances_
indices=np.argsort(importances)
```

In [27]:

```
plt.title('Feature Importances')
plt.barh(range(len(indices)),importances[indices],color='b',align='center')
plt.yticks(range(len(indices)),[features[i]for i in indices])
plt.xlabel('Realative Importances')
plt.show()
```



In [29]:

```
del data['AgeCategory']
del data['Workclass']
del data['Education']
del data['EducationNum']
del data['Occupation']
del data['Sex']
del data['FnlwgtCategory']
del data['Race']
del data['CapitalLossCategory']
del data['HoursPerWeekCategory']
```

In [30]:

```
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
```

In [31]:

```
(x_train,x_test,y_train,y_test)=train_test_split(all_input,all_class,train_size=0.67,random_state=10)
```

In [32]:

```
clf = DecisionTreeClassifier(random_state=10)
clf.fit(x_train,y_train)
```

Out[32]:

```
DecisionTreeClassifier
DecisionTreeClassifier(random_state=10)
```

In [33]:

```
y_train_pred=clf.predict(x_train)
y_test_pred=clf.predict(x_test)
```

In [34]:

```
print(y_train_pred)
y_train
```

```
[0 1 0 ... 0 0 0]
```

Out[34]:

```
39203    0
16702     1
43825     0
48735     1
34480     0
..
40059     1
28017     0
29199     0
40061     0
17673     0
Name: Class, Length: 32724, dtype: int64
```

In [35]:

```
from sklearn import metrics,model_selection,preprocessing
wrong_train_pred=(y_train !=y_train_pred).sum()
print("Total wrong detected on training data= {}".format(wrong_train_pred))

accuracy_train=metrics.accuracy_score(y_train,y_train_pred)
print("Accuracy of this model on training data= {:.3f}".format(accuracy_train))
```

```
Total wrong detected on training data= 2812
Accuracy of this model on training data= 0.914
```

In [37]:

```
wrong_test_pred=(y_test !=y_test_pred).sum()
print("Total wrong detected on test data = {}".format(wrong_test_pred))

accuracy_test=metrics.accuracy_score(y_test,y_test_pred)
print("Accuracy of this model on test data = {:.3f}".format(accuracy_test))
```

```
Total wrong detected on test data = 3058
Accuracy of this model on test data = 0.810
```

In [38]:

```
train_accuracy=[]
test_accuracy=[]
train_error=[]
valid_error=[]
test_error=[]
for depth in range(1,40):
    dt_model_tree=DecisionTreeClassifier(max_depth=depth,random_state=10)
    dt_model_tree.fit(x_train,y_train)
    train_accuracy.append(dt_model_tree.score(x_train,y_train))
    test_accuracy.append(dt_model_tree.score(x_test,y_test))
```

In [39]:

```
frame = pd.DataFrame({'max_depth': range(1,40), 'train_acc':train_accuracy, 'test_acc':test_accuracy})
frame.head()
```

Out[39]:

| | max_depth | train_acc | test_acc |
|---|-----------|-----------|----------|
| 0 | 1 | 0.761062 | 0.760020 |
| 1 | 2 | 0.814876 | 0.812446 |
| 2 | 3 | 0.821691 | 0.819705 |
| 3 | 4 | 0.831653 | 0.828577 |
| 4 | 5 | 0.832508 | 0.829445 |

In [40]:

```
import numpy as np

train_accuracy = np.array(train_accuracy)
train_error = (1 - train_accuracy) * 32562

test_accuracy = np.array(test_accuracy)
test_error = (1 - test_accuracy) * 8140
```

In [41]:


```
from IPython.display import Image,display
import matplotlib.pyplot as plt,pydotplus
import graphviz
```

In [42]:

```
import matplotlib.pyplot as plt,pydotplus
plt.figure(figsize=(12,6))
plt.plot(frame['max_depth'],frame['train_acc'],label='Training',marker='o')
plt.plot(frame['max_depth'],frame['test_acc'],label='Test',marker='v')
plt.xlabel('Depth of tree')
plt.ylabel('Perfromance accuracy')
plt.legend()
```

Out[42]:

<matplotlib.legend.Legend at 0x288d5530580>



In [43]:


```
frame1 = pd.DataFrame({'max_depth': range(1,40),'train_err':train_error,'test_err':test_error})
```

In [44]:

```
import matplotlib.pyplot as plt,pydotplus
plt.figure(figsize=(12,6))
plt.plot(frame1['max_depth'],frame1['train_err'],label='Training',marker='o')
plt.plot(frame1['max_depth'],frame1['test_err'],label='Test',marker='v')
plt.xlabel('log(Vertices)(as a fuction of depth of tree i.e. vertices = 2^(n+1) - 1)')
plt.ylabel('Number of Error')
plt.legend()
```

Out[44]:

<matplotlib.legend.Legend at 0x288d5594760>



In [49]:

```
from sklearn import tree
plt.figure(figsize=(200, 20))
tree.plot_tree(clf, filled=True,max_depth=5, feature_names=['HoursPerWeekCategory','CapitalLossCategory','CapitalGainCategory'])
```

Text(0.5785123966942148, 0.21428571428571427, 'EducationNum <= 2.5\ngini = 0.353\nsamples = 140\nvalue = [32, 108]\nnclass = No'),
Text(0.5702479338842975, 0.07142857142857142, '\n (...) \n'),
Text(0.5867768595041323, 0.07142857142857142, '\n (...) \n'),
Text(0.6115702479338843, 0.21428571428571427, 'HoursPerWeekCategory <= 13.5\ngini = 0.499\nsamples = 187\nvalue = [90, 97]\nnclass = No'),
Text(0.6033057851239669, 0.07142857142857142, '\n (...) \n'),
Text(0.6198347107438017, 0.07142857142857142, '\n (...) \n'),
Text(0.6942148760330579, 0.5, 'HoursPerWeekCategory <= 12.5\ngini = 0.095\nsamples = 17250\nvalue = [1639, 858]\nnclass = Yes'),
Text(0.6611570247933884, 0.35714285714285715, 'HoursPerWeekCategory <= 11.5\ngini = 0.072\nsamples = 1632\nvalue = [15716, 610]\nnclass = Yes'),
Text(0.6446280991735537, 0.21428571428571427, 'AgeCategory <= 1.5\ngini = 0.048\nsamples = 13824\nvalue = [13487, 337]\nnclass = Yes'),
Text(0.6363636363636364, 0.07142857142857142, '\n (...) \n'),
Text(0.6528925619834711, 0.07142857142857142, '\n (...) \n'),
Text(0.6776859504132231, 0.21428571428571427, 'AgeCategory <= 1.5\ngini = 0.194\nsamples = 2502\nvalue = [2229, 273]\nnclass = Yes'),
Text(0.6694214876033058, 0.07142857142857142, '\n (...) \n'),
Text(0.6859504132231405, 0.07142857142857142, '\n (...) \n'),

In []: